

Unit 6: Radial-Basis Function Networks

A Hybrid Approach to Pattern Classification

Kiran Bagale

August, 2025

Introduction

Traditional vs. Hybrid Approach

Traditional Approach

- Back-propagation learning
- Stochastic approximation
- Single-stage training

Hybrid Approach (RBF)

- Two-stage procedure
- Nonlinear transformation
- Linear classification

Two-Stage Hybrid Approach

Stage 1: Nonlinear Transformation

Objective

Transform nonlinearly separable patterns into a new set where patterns are likely to become linearly separable

- Mathematical foundation: Cover's theorem (1965)
- Increases likelihood of linear separability
- Unsupervised training of hidden layer


Stage 2: Linear Classification

Objective

Complete the classification using least-squares estimation

- Uses transformed feature space
- Linear output layer
- Supervised training

RBF Network Architecture: Three-Layer Structure



Layer Functions

1 Input Layer

- Source nodes (sensory units)
- Connects network to environment

2 Hidden Layer

- Applies nonlinear transformation
- High dimensionality
- Trained unsupervised (Stage 1)

3 Output Layer

- Linear transformation
- Supplies network response
- Trained supervised (Stage 2)

Cover's Theorem Requirements

Two Conditions for Linear Separability

- 1 Nonlinear transformation from input space to hidden space
- 2 High dimensionality of the hidden space

Key Insight

RBF networks satisfy both conditions of Cover's theorem, making linear separability in the feature space highly likely.

Gaussian Radial-Basis Function

- Most important member of RBF class
- Mathematical form: $\phi(\mathbf{x}) = \exp\left(-\frac{\|\mathbf{x}-\mathbf{c}\|^2}{2\sigma^2}\right)$
- Can be viewed as a **kernel**

Kernel Method

Two-stage procedure based on Gaussian function:

- Kernel transformation (Stage 1)
- Linear classification (Stage 2)

Connection to Statistics

- RBF networks relate to kernel regression in statistics
- Bridge between neural networks and statistical methods
- Theoretical foundation for understanding RBF performance

Neural Networks \leftrightarrow Statistical Methods

Advantages of RBF Networks


- Faster training than backpropagation
- Strong theoretical foundation
- Universal approximation capability
- Good generalization properties
- Clear separation of learning stages

6.2 Cover's Theorem on the Separability of Patterns

The Fundamental Challenge

Real-world classification problems:

- Often nonlinearly separable
- Linear classifiers fail
- Need complex decision boundaries



Cover's Key Insight (1965)

Nonlinearly separable problems in low dimensions become linearly separable in high dimensions with high probability!

Problem Setup

Given

- Input patterns: $\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_N \in \mathbb{R}^n$
- Class labels: $y_i \in \{-1, +1\}$
- Nonlinear mapping: $\phi : \mathbb{R}^n \rightarrow \mathbb{R}^m$ where $m \gg n$

Linear Separability in Feature Space

Exists hyperplane $\mathbf{w}^T \phi(\mathbf{x}) + b = 0$ such that:

$$\mathbf{w}^T \phi(\mathbf{x}_i) + b > 0 \quad \text{for class } +1 \quad (1)$$


$$\mathbf{w}^T \phi(\mathbf{x}_i) + b < 0 \quad \text{for class } -1 \quad (2)$$

Mathematical Intuition

In high dimensions, the volume grows exponentially, providing exponentially more "room" for separating hyperplanes.


The Transformation Process

Nonlinearly Separable



Input Space \mathbb{R}^n

Linearly Separable



Feature Space \mathbb{R}^m

$$\phi(\mathbf{x})$$





Figure: Separability of Patterns

Theorem Statement

Theorem (Cover's Theorem)

Given N patterns $\{\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_N\}$ in \mathbb{R}^n , if mapped nonlinearly into feature space of dimension m , the probability of linear separability is:

$$P(\text{linear separability}) = \frac{1}{2^N} \sum_{k=0}^{m-1} \binom{N-1}{k}$$

Key Conditions

- 1 General nonlinear mapping (not data-specific)
- 2 Sufficiently large feature dimension m
- 3 Patterns in "general position"

Critical Cases

Case 1: $m \geq N$

$$P = 1$$

Perfect separability!

High dimensions provide enough "room" for any separating hyperplane.

Case 2: $m < N$

$$P = \frac{1}{2^{N-1}} \sum_{k=0}^{m-1} \binom{N-1}{k}$$

Probability increases exponentially as $m \rightarrow N$.


Key Insight

As feature dimension increases toward the number of patterns, linear separability becomes almost certain!

Example: N=4 Patterns

For $N = 4$ patterns: $P = \frac{1}{8} \sum_{k=0}^{m-1} \binom{3}{k}$

m	Binomial Sum	P	Percentage
1	$\binom{3}{0} = 1$	$\frac{1}{8}$	12.5%
2	$1 + \binom{3}{1} = 4$	$\frac{4}{8}$	50%
3	$1 + 3 + \binom{3}{2} = 7$	$\frac{7}{8}$	87.5%
≥ 4	$1 + 3 + 3 + 1 = 8$	$\frac{8}{8}$	100%



XOR Problem Transformation

Input Space (\mathbb{R}^2): Not linearly separable

Input \mathbf{x}	Class	2D Coordinates
(0, 0)	-1	(0, 0)
(0, 1)	+1	(0, 1)
(1, 0)	+1	(1, 0)
(1, 1)	-1	(1, 1)

Feature Space (\mathbb{R}^3): $\phi(\mathbf{x}) = [x_1, x_2, x_1x_2]^T$


Input	Feature Vector	Class
(0, 0)	$[0, 0, 0]^T$	-1
(0, 1)	$[0, 1, 0]^T$	+1
(1, 0)	$[1, 0, 0]^T$	+1
(1, 1)	$[1, 1, 1]^T$	-1

Separating hyperplane: $w_1x_1 + w_2x_2 - 2w_3x_1x_2 + b = 0$

Curse vs. Blessing of Dimensionality

Low Dimensions

- Points are "crowded"
- Limited hyperplane orientations
- Complex boundaries needed



High Dimensions

- Points spread out
- Many hyperplane orientations
- Linear separation likely






Figure: Three examples of φ -separable dichotomies of different sets of five points in two dimensions: (a) linearly separable dichotomy; (b) spherically separable dichotomy; (c) quadratically separable dichotomy.

Asymptotic Behavior

For large N and $m \approx N$:

$$P \approx \frac{1}{\sqrt{\pi N/2}} \sum_{k=0}^{m-1} \exp\left(-\frac{(k - N/2)^2}{N/2}\right)$$



Kernel Interpretation

The nonlinear mapping $\phi(\mathbf{x})$ defines a kernel function:

$$K(\mathbf{x}, \mathbf{x}') = \phi(\mathbf{x})^T \phi(\mathbf{x}')$$

Popular Kernels

- **Gaussian RBF:** $K(\mathbf{x}, \mathbf{x}') = \exp\left(-\frac{\|\mathbf{x}-\mathbf{x}'\|^2}{2\sigma^2}\right)$
- **Polynomial:** $K(\mathbf{x}, \mathbf{x}') = (\mathbf{x}^T \mathbf{x}' + 1)^d$
- **Sigmoid:** $K(\mathbf{x}, \mathbf{x}') = \tanh(\alpha \mathbf{x}^T \mathbf{x}' + \beta)$

Key Advantage

We can work with kernels without explicitly computing $\phi(\mathbf{x})$ - the "kernel trick"!

Reproducing Kernel Hilbert Spaces (RKHS)

Cover's theorem explains why kernel methods work:

- 1 Kernel implicitly maps to high-dimensional space
- 2 Linear separability highly probable in this space
- 3 Optimization becomes convex (linear in feature space)
- 4 Strong theoretical guarantees



RBF Network Design Principles

Principle 1: High-Dimensional Hidden Layer

Use $m \gg n$ hidden units to ensure high probability of linear separability

Principle 2: Appropriate Basis Functions

Choose RBF centers \mathbf{c}_i and widths σ_i :

$$\phi_i(\mathbf{x}) = \exp\left(-\frac{\|\mathbf{x} - \mathbf{c}_i\|^2}{2\sigma_i^2}\right)$$

Principle 3: Two-Stage Learning

- ① **Stage 1:** Learn \mathbf{c}_i, σ_i (unsupervised)
- ② **Stage 2:** Learn output weights (supervised, linear)

Trade-offs

Benefits

- High probability of separability
- Simple linear learning stage
- Strong theoretical foundation
- Universal approximation

Costs

- Computational complexity
- Risk of overfitting
- Need more training data
- Parameter selection

Design Guideline

Choose m large enough for separability but not so large as to cause overfitting: typically $m = O(\sqrt{N})$ to $O(N)$.

Assumptions and Limitations

Key Assumptions

- **General Position:** Patterns not in special geometric arrangements
- **Random Mapping:** Transformation not specifically designed for data
- **Infinite Precision:** Mathematical idealization

Practical Considerations

- **Finite Sample Effects:** Real datasets have limited samples
- **Noise Sensitivity:** High dimensions can amplify noise
- **Computational Reality:** Memory and time constraints

Bottom Line


Cover's theorem provides theoretical justification, but practical implementation requires careful parameter tuning and regularization.

Key Takeaways

- ① **Core Insight:** Nonlinear transformation to high dimensions makes linear separability highly probable
- ② **Mathematical Foundation:** Probability approaches 1 as $m \rightarrow N$
- ③ **Practical Impact:** Justifies two-stage learning in RBF networks
- ④ **Kernel Connection:** Explains success of kernel methods and SVMs
- ⑤ **Design Guidance:** Choose sufficient but not excessive dimensionality

Cover's theorem bridges theory and practice in pattern recognition

The Big Picture



Unifying theoretical framework for nonlinear classification

6.3 The Interpolation Problem: From Theory to Practice

Cover's Theorem Application

Cover's theorem shows that non-linear mapping to high dimensions enables linear separability. But how do we implement this practically?

The Solution:

- Transform nonlinear classification into interpolation
- Use radial-basis functions
- Two-stage hybrid approach

Two-Phase Learning Process

Training Phase

Optimization of a fitting procedure for surface Γ , based on known input-output examples (patterns).

Generalization Phase

Synonymous with **interpolation** between data points, with interpolation performed along the constrained surface generated by the fitting procedure.


The Network as a Hypersurface

Consider a feedforward network: Input \rightarrow Hidden \rightarrow Output

Mathematical Representation

The network represents a mapping: $s : \mathbb{R}^{m_0} \rightarrow \mathbb{R}^1$

We can think of this map s as a **hypersurface** $\Gamma \subset \mathbb{R}^{m_0+1}$



The surface Γ is unknown and training data are contaminated with noise. **This leads us to the theory of multivariable interpolation in high dimensional space, with a long history (Davis, 1963).**

The Strict Interpolation Problem

Problem Statement

Given a set of N different points $\{\mathbf{c}_i \in \mathbb{R}^{m_0} | i = 1, 2, \dots, N\}$ and corresponding real numbers $\{d_i \in \mathbb{R}^1 | i = 1, 2, \dots, N\}$, find a function $F : \mathbb{R}^{m_0} \rightarrow \mathbb{R}^1$ that satisfies the interpolation condition:

$$F(\mathbf{c}_i) = d_i, \quad i = 1, 2, \dots, N$$

Strict Interpolation

The interpolating surface (function F) is constrained to pass through **all** the training data points.


Radial-Basis Functions Solution

The RBF technique chooses a function F of the form:

$$F(\mathbf{x}) = \sum_{i=1}^N w_i \varphi(\|\mathbf{x} - \mathbf{c}_i\|)$$

where:

- $\{\varphi(\|\mathbf{x} - \mathbf{c}_i\|) | i = 1, 2, \dots, N\}$ are **radial-basis functions**
- $\|\cdot\|$ denotes a norm (usually Euclidean)
- $\mathbf{c}_i \in \mathbb{R}^{m_0}$ are the **centers** of the RBFs
- w_i are the unknown weights to be determined



Linear System of Equations

Inserting the interpolation conditions into the RBF expansion:

$$\sum_{j=1}^N w_j \varphi(\|\mathbf{x}_i - \mathbf{c}_j\|) = d_i, \quad i = 1, 2, \dots, N$$

This gives us the matrix equation:

$$\begin{bmatrix} \varphi_{11} & \varphi_{12} & \cdots & \varphi_{1N} \\ \varphi_{21} & \varphi_{22} & \cdots & \varphi_{2N} \\ \vdots & \vdots & \ddots & \vdots \\ \varphi_{N1} & \varphi_{N2} & \cdots & \varphi_{NN} \end{bmatrix} \begin{bmatrix} w_1 \\ w_2 \\ \vdots \\ w_N \end{bmatrix} = \begin{bmatrix} d_1 \\ d_2 \\ \vdots \\ d_N \end{bmatrix}$$

where $\varphi_{ij} = \varphi(\|\mathbf{x}_i - \mathbf{c}_j\|)$

Compact Matrix Form

Let:

$$\mathbf{d} = [d_1, d_2, \dots, d_N]^T \quad (\text{desired response vector}) \quad (3)$$

$$\mathbf{w} = [w_1, w_2, \dots, w_N]^T \quad (\text{linear weight vector}) \quad (4)$$

$$\Phi = \{\varphi_{ij}\}_{i,j=1}^N \quad (\text{interpolation matrix}) \quad (5)$$

The system becomes:

$$\Phi \mathbf{w} = \mathbf{d}$$

Solution

Assuming Φ is nonsingular (invertible):

$$\mathbf{w} = \Phi^{-1} \mathbf{d}$$

Critical Question

How can we ensure the interpolation matrix Φ is nonsingular?

Micchelli's Theorem

Theorem (Micchelli, 1986)

Let $\{\mathbf{x}_i\}_{i=1}^N$ be a set of distinct points in \mathbb{R}^{m_0} . Then the N -by- N interpolation matrix Φ , whose (i,j) -th element is $\varphi_{ij} = \varphi(\|\mathbf{x}_i - \mathbf{c}_j\|)$, is nonsingular.

Key Requirement

The points $\{\mathbf{c}_i\}_{i=1}^N$ must all be **distinct** (different from each other).

This theorem covers a large class of radial-basis functions of particular interest in RBF networks.

Important RBF Functions

Micchelli's theorem applies to:

1. Multiquadrics (Hardy, 1971)

$$\varphi(r) = (r^2 + c^2)^{1/2} \quad \text{for some } c > 0, r \in \mathbb{R}$$

2. Inverse Multiquadrics (Hardy, 1971)

$$\varphi(r) = \frac{1}{(r^2 + c^2)^{1/2}} \quad \text{for some } c > 0, r \in \mathbb{R}$$

3. Gaussian Functions


$$\varphi(r) = \exp\left(-\frac{r^2}{2\sigma^2}\right) \quad \text{for some } \sigma > 0, r \in \mathbb{R}$$

All require only that data points be **distinct** - no constraints on data size N or dimensionality m_0 .

Properties of RBF Functions


Localized Functions:

- Inverse multiquadrics
- Gaussian functions
- Property: $\varphi(r) \rightarrow 0$ as $r \rightarrow \infty$



Growing Functions:

- Multiquadrics
- Property: $\varphi(r) \rightarrow \infty$ as $r \rightarrow \infty$
- Still yield nonsingular matrices!



Remarkable Result

Functions that grow to infinity can still be used for smooth input-output mapping with great accuracy, yielding positive-definite interpolation matrices.

Matrix Properties and Implications

For Localized Functions (Inverse Multiquadrics, Gaussian)

- Interpolation matrix Φ is positive definite
- All eigenvalues are positive
- Matrix is well-conditioned for numerical computation

For Growing Functions (Multiquadrics)

- Matrix has $N - 1$ negative eigenvalues
- Only one positive eigenvalue
- Matrix is NOT positive definite
- But still nonsingular and suitable for RBF networks!

Practical Implication

Both localized and growing RBF functions can achieve smooth input-output mapping with great accuracy (Powell, 1988).

Summary: Interpolation Problem

- ➊ **Problem Transformation:** Neural network learning becomes multivariable interpolation
- ➋ **RBF Solution:** Use radial basis functions with centers at data points
- ➌ **Linear System:** Results in solvable linear equation $\Phi \mathbf{w} = \mathbf{d}$
- ➍ **Theoretical Guarantee:** Micchelli's theorem ensures matrix nonsingularity
- ➎ **Function Classes:** Multiquadrics, inverse multiquadrics, and Gaussian functions all work
- ➏ **Practical Success:** Both localized and growing functions achieve high accuracy

The interpolation approach provides a solid mathematical foundation for RBF network design

6.4 Radial-Basis-Function Networks

Overview

In light of interpolation theory, we can envision a **radial-basis-function (RBF)** network as a layered structure with three layers:

- 1 **Input layer:** consists of m_0 source nodes, where m_0 is the dimensionality of the input vector \mathbf{x}
- 2 **Hidden layer:** consists of the same number of computation units as the size of the training sample, namely N
- 3 **Output layer:** consists of a single computational unit in the RBF structure

RBF Network Structure Each unit in the hidden layer is mathematically described by a radial-basis function:

$$\varphi_j(\mathbf{x}) = \varphi(\|\mathbf{x} - \mathbf{c}_j\|), \quad j = 1, 2, \dots, N$$

Key characteristics:

- The j th input data point \mathbf{c}_j defines the center of the radial-basis function
- The vector \mathbf{x} is the signal (pattern) applied to the input layer
- Links connecting source nodes to hidden units are **direct connections with no weights**
- Output layer size is typically much smaller than the hidden layer

Gaussian Radial-Basis Function

Mathematical Formulation

We focus on using a **Gaussian function** as the radial-basis function:

$$\begin{aligned}\varphi_j(\mathbf{x}) &= \varphi(\mathbf{x} - \mathbf{c}_j) \\ &= \exp\left(-\frac{1}{2\sigma_j^2}\|\mathbf{x} - \mathbf{c}_j\|^2\right), \quad j = 1, 2, \dots, N \quad (5.20)\end{aligned}$$

where:

- σ_j is a measure of the **width** of the j th Gaussian function with center \mathbf{c}_j
- Typically, all Gaussian hidden units are assigned a common width σ
- The parameter that distinguishes one hidden unit from another is the center \mathbf{c}_j

Practical Modifications to RBF Networks: Addressing Real-World Issues

Challenge: The theoretical RBF network formulation has practical limitations:

- Training samples $\{\mathbf{x}_i, d_i\}_{i=1}^N$ are typically **noisy**
- Interpolation based on noisy data can lead to misleading results
- Having a hidden layer of the same size as the input layer is computationally wasteful for large training samples

Theoretical vs Practical RBF Networks

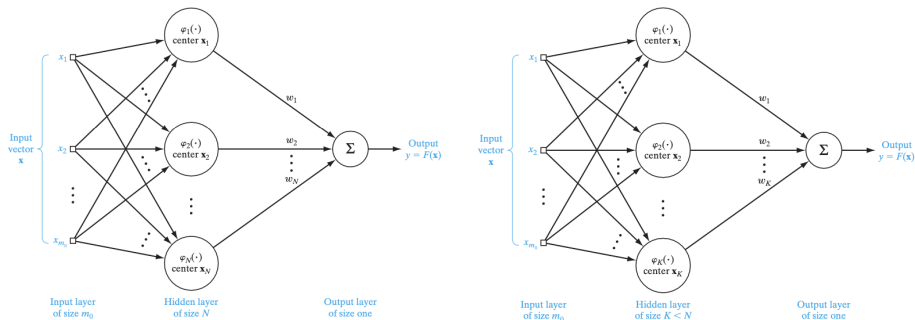


Figure: Structure of a practical RBF networks

Note: $x_j = c_j$

Improved RBF Network Structure

Reducing Computational Complexity

Key insight: Correlation between adjacent data points in the training sample creates redundancy in the hidden layer.

Practical design principle: Make the size of the hidden layer a fraction of the size of the training sample.

Advantages

- Reduces computational resources
- Networks in Fig. (a) and (b) share a common mathematical structure
- Unlike multilayer perceptron, RBF training does **not involve back propagation** of error signals

Mathematical Form of Practical RBF Networks

General Approximating Function

The approximating function realized by practical RBF structures:

$$F(\mathbf{x}) = \sum_{j=1}^K w_j \varphi(\mathbf{x}, \mathbf{c}_j)$$

where:

- Input vector dimensionality: m_0
- Each hidden unit characterized by: $\varphi(\mathbf{x}, \mathbf{c}_j)$, $j = 1, 2, \dots, K$
- $K < N$ (smaller than training set size)
- Output layer characterized by weight vector \mathbf{w} with dimensionality K

Comparison: Theoretical vs. Practical RBF Networks

Key Differences

Theoretical RBF (Fig. (a))

- Hidden layer dimensionality: N
- N = size of training set
- Centers defined by input vectors
- Noiseless training assumption

Practical RBF (Fig. (b))

- Hidden layer dimensionality:
 $K < N$
- Reduced computational complexity
- Requires new procedure for center selection
- Handles noisy training data

Next Steps

The next section addresses center selection procedures for practical RBF networks using Gaussian functions.

RBF Network Training Algorithm

Input

- Training data: $\{(\mathbf{x}_i, y_i)\}$ for $i = 1$ to N
- Number of RBF centers: K (randomly selected)

Algorithm Steps

1 Choose centers (\mathbf{c}_j):

Use K-Means clustering or randomly select K samples from the training data.

2 Choose spread (σ):

A common choice: $\sigma = \frac{d_{\max}}{\sqrt{2K}}$

where d_{\max} is the maximum distance between any two centers.

RBF Network Algorithm (continued)

Algorithm Steps (continued)

3 Compute RBF activations for each input:

For each input \mathbf{x}_i and center \mathbf{c}_j , compute:

$$\Phi_{ij} = \varphi_j(\mathbf{x}_i) = \exp\left(-\frac{\|\mathbf{x}_i - \mathbf{c}_j\|^2}{2\sigma^2}\right) \quad (6)$$

4 Train output weights using linear regression:

$$\mathbf{w} = (\Phi^T \Phi)^{-1} \Phi^T \mathbf{y} \quad (7)$$

5 Predict for new input \mathbf{x} :

$$\hat{y} = \sum_{j=1}^K w_j \cdot \varphi_j(\mathbf{x}) + b \quad (8)$$

Thank You!!!